

LogitTrail: An Inspectable Local Interface for Typed Decisions from Frozen Language Models

Yuki Oshio
UPHASH Inc.
oshio@uphash.net

Abstract

Many language-model applications need a small typed decision rather than free-form text. LogitTrail maps a state and criteria to a categorical choice, a true-candidate probability, or an expected ordinal stage using frozen next-token logits, with no trained decision head. Its local browser exposes typed values and distributions; a separate model-free explorer aligns saved conditions by candidate meaning and traces results to source records and physical calls. A 4,050-request matched comparison finds identical direct-readout and native one-token outputs in all 1,350 pairs, without an established latency advantage. Two subsequent studies across three checkpoints use 7,600 and 14,400 requests to examine presentation sensitivity and probability averaging. Reduced sensitivity can coexist with incorrect or nearly constant predictions, while averaging costs multiple model calls. The contribution is a runnable inspection workflow and auditable system evaluation separating type validity, candidate concentration, task quality, and computation cost.

1 Introduction

Applications may need a route, an urgency score, or a rubric stage. Typed interfaces make these contracts explicit, but valid types and concentrated candidate probabilities do not establish correct decisions.

LogitTrail (formerly Mini Jev) demonstrates this distinction through a local browser interface, an HTTP/Python API, and a saved-evidence explorer. Users can run typed questions locally or inspect recorded comparisons without downloading a model. The interface is inspired by TypeSafe’s Jev (Almeida, 2026; TypeSafe, 2026b); this independent implementation does not reproduce its architecture, training, calibration, compatibility, or performance.

Developers need to determine what changed in a decision, whether the change improved it, and which executions support that conclusion. Our workflow connects semantic candidate comparisons to reference outcomes, panel frequencies, and source provenance. Its mechanism—frozen candidate scoring, restricted softmax, and typed response construction—is conventional (Schick and Schütze, 2021; Ma et al., 2025; LMQL contributors, 2026). Neither scoring nor typed probability interfaces are new.

Schema validity, task quality, and paired timing answer different questions. Our matched study does not establish a latency advantage over native one-token selection. Source and artifacts: github.com/UpHash-Network/mini-jev. Browser access: [saved-evidence explorer](#). Video: [captioned demonstration](#).

2 System and Typed Contract

2.1 From a state to typed values

A request contains a shared state and named questions. Each question supplies a type, natural-language instructions, and candidate criteria. The service returns a semantic label, a candidate distribution, and a type-specific value:

- **Choice** returns the key of the most likely candidate. Keys are canonically sorted before prompting.
- **Noul** exposes the probability of the true candidate in a two-candidate false/true decision.
- **Score** returns the expected zero-based stage of an ordered rubric; the most likely stage is also available as the label.

The distinction between Score’s expectation and its hard label matters. An expectation of 1.6 on stages 0–3 is a probability-weighted position,

not an additional model-generated stage or a calibrated utility. Its interpretation assumes equally spaced stage indices; users must decide whether that fits their rubric.

Let z_i denote the next-token logit for candidate $i \in \{0, \dots, K-1\}$, with K candidates and temperature $T > 0$. The service computes

$$p_i = \frac{\exp(z_i/T)}{\sum_{j=0}^{K-1} \exp(z_j/T)},$$

$$s = \sum_{i=0}^{K-1} i p_i. \quad (1)$$

These are probabilities conditional on the selected candidate set. They are neither probabilities over all valid textual answers nor guarantees that the criteria are exhaustive. Changing the criteria changes the question as well as the normalization. The optional scalar temperature preserves the maximum-logit label, but can change Noul and Score values.

The interface also displays entropy concentration,

$$c = 1 - H(p)/\log K,$$

$$H(p) = - \sum_i p_i \log p_i. \quad (2)$$

A concentrated distribution has c near one. The display identifies this quantity as concentration rather than estimated accuracy. The response includes probability-semantics and calibration metadata, so downstream applications can retain the same distinction.

2.2 Local execution

The released inference configuration uses Qwen3.6-35B-A3B in Q4_K_M GGUF form (Qwen Team, 2026; ggml-org, 2026), through a source-built llama.cpp helper (ggml-org and contributors, 2026). The model and its output projection remain frozen; there is no trained decision head. The helper performs the full vocabulary projection and then gathers candidate logits. It does not implement a selected-row projection optimization.

The prompt repeats the task twice, following the use of repetition for non-reasoning inference (Leviathan et al., 2025). Choice and Noul use letter labels; Score with at most ten stages uses digits after a prefilled answer-object prefix. Larger Score rubrics use letters. The implementation checks that every candidate is a distinct single token at the actual prefix boundary. The API accepts 2–26 candidates; quality evaluation covers

2–8, while larger candidate sets have functional coverage only.

A multi-question request executes its questions sequentially. Model state is cleared between questions: they share the supplied state text, but not a cached prefill. Complete inputs exceeding 2,048 tokens are rejected rather than silently truncated. Schema validation and deterministic response construction prevent free-form output from crossing the API boundary; these software properties do not prevent semantic errors.

2.3 Relationship to one-token decoding

Restricted softmax readout has the same mathematical candidate distribution as a sampler that masks all non-candidates, provided the model prefix, logits, candidate IDs, temperature, and other processors match. Greedy labels then agree unless tie-breaking differs; stochastic sampling agrees in distribution, not necessarily in realized samples. This is an equivalence under specified conditions, not a new inference theorem or learning method.

In particular, selecting a first token after prefill need not require another model forward pass. Avoiding a generated token therefore does not itself establish a speedup. The matched study in Section 4 tests this point using an actual native one-token sampler, rather than comparing direct readout only with a long textual answer.

3 Demonstration and Access

3.1 Inspect saved evidence without a model

The static [saved-evidence explorer](#) requires only a browser. It contains all eligible records from the presentation and averaging studies: 22,000 physical calls and 9,000 derived answers in 18 panels. Their 3,600 study/item/model combinations represent 1,200 distinct source items, reused across checkpoints. These are recorded observations, not new inference.

Users choose a study, checkpoint, task, item, and two conditions. Both sides retain the same item and model; probabilities align by semantic key on a common 0–100% scale even when token bindings or row order differ. The display separates the selected label, typed value, concentration, and reference. Full-panel label frequencies remain visible independently of item selection. This permits inspection of an order flip, a correct answer lost after averaging, or a nearly constant predictor. Four outcome-selected examples illustrate these cases;

they do not estimate prevalence.

The explorer also distinguishes executions from derived results. An averaged answer lists its physical members; comparisons count their union without double-counting shared calls. Source details identify archive members, JSONL lines, mappings, and hashes. Loaded panels are SHA-256 checked; comparison exports retain exact values and sources. Original task text is withheld. URL fragments share selections. A functional audit checks 172,200 dependent within-item pairs, 31,000 condition/source rows, 12,600 Score conditions, and 155 panel distributions against raw-archive calculations, with no discrepancy. A linked complete-JSON view retains the same information. These are software checks, not human-performance results.

3.2 Run local typed decisions

The separate live browser connects through a Python bridge to the local model API. Users edit a shared state, instructions, and criteria for Choice, Noul, and Score, then select *Run decisions*. Results expose candidate probabilities, concentration, elapsed time, and request/response JSON. Editing marks prior results stale. The live UI has no replay or cyclic-averaging mode: an unavailable API produces an error. Input edits explore behavior, not controlled accuracy.

A linked [reviewer guide](#) proceeds from the explorer to model-free verification and optional live installation. The tested live environment is a 64 GB Apple M5 Pro; its model file is approximately 20.4 GB. Appendix A gives installation and verification boundaries. The static explorer has no model, API-key, or build requirement.

4 Evidence Behind the Demonstration

4.1 Evaluation panels and provenance

Table 1 separates the evidence sources. Repeated requests are not additional quality examples, and the external tasks do not share one accuracy metric.

The historical local suite contains 2,220 generated items from 45 template families and 180 individually AI-authored items; its legacy manual field does not mean human annotation. Authoring used earlier errors, so this is regression material rather than a blind test. All 2,400 outputs were valid, but 162 top labels disagreed with authored reference labels: structural validity does not estab-

Evidence	Size	Interpretation
Local regression	2,400 items	Self-authored Japanese; 800 per type
Expanded external	600 items	JCoLA, JSTS, JCQA; 200 each
Matched run-time	4,050 quests	re- 150 local $\times$ 5 repetitions, plus 600 external; 3 modes
Presentation study	7,600 quests	re- Same 600 external items; 3 checkpoints
Cyclic averaging	14,400 requests	600 additional external items; 3 checkpoints

Table 1: Evaluation accounting. The last two studies use 1,200 unique source items in total. Measured counts exclude synthetic checks and warm-ups. JCQA abbreviates JCommonsenseQA.

lish correctness. Temperature fitting did not improve all calibration and score-error metrics.

The expanded sample has 200 public-development examples each from JCoLA (Someya et al., 2024), JSTS, and JCommonsenseQA (Kurihara et al., 2022). Hash selection and rubrics were frozen before inference, without independent preregistration. Six JSTS rows share text or image provenance with an earlier 300-item JNLI pilot. Public development data may have appeared in model training; these are not independent hidden-test evaluations.

4.2 Matched runtime and parity

All three conditions used one resident model and a common research HTTP response containing a type and semantic label. Direct readout and native one-token greedy selection used identical complete prompts and candidates. The JSON condition actually generated a grammar-constrained object, with a serialization-specific prompt. Its timing comparison therefore includes prompt and pre-fill differences as well as multi-token generation; it does not isolate sampler overhead.

The schedule contained $(150 \times 5 + 600) \times 3 = 4,050$ measured requests, plus 21 synthetic warm-ups. All completed with valid responses and no recorded HTTP failures. Mode order was randomized within each item/repetition. Timing covered serialization through loopback HTTP parsing, including native audit transfer but excluding trace-file writes. These are research-harness timings, not browser rendering times.

Direct and one-token conditions matched prompts, token sequences, candidate IDs, logits, and labels in all 1,350 paired observations. Maximum observed candidate-logit difference was

Candidate comparison

[Open complete JSON for this selection](#) - Includes the selected comparison, full panel, and source index.

A Recorded condition

Original presentation

The saved original presentation.

B Recorded condition

Reversed candidate rows

Reverse only the displayed rows, preserving answer-token, key, and meaning bindings.

The most likely answer stays the same. Largest candidate probability change: 7.384769 percentage points.

Most likely semantic answer

A true
B true

Recorded tie handling is preserved.

True-candidate probability

A 0.986484
B 0.912636

Within the false/true candidate pair.

Concentration

A 0.8967
B 0.5724

0 = uniform - 1 = peaked; not accuracy.

Model calls per result

A 1
B 1

0 shared; 2 distinct calls across A and B.

Probability within the candidate set

A and B share a 0–100% scale. Δ is B – A in percentage points. Values are aligned by semantic key, even when presentation order changes.

false · unacceptable
key: false

A
B



1.351581%
8.73635%
Δ percentage points

true · acceptable
key: true

A
B



98.648419%
91.26365%
Δ percentage points

Dataset reference label: true. A matches; B matches. This is one recorded item, not a new accuracy estimate.
► Answer frequencies across this full panel

Concentration measures how peaked a candidate distribution is; it is not the probability of being correct. Score expectations assume equally spaced stages. Recorded examples do not establish human usability or a causal model-size effect.

Figure 1: The saved-evidence explorer aligns recorded candidate probabilities by meaning. Users inspect labels, typed values, references, and physical-call costs, then trace or export the comparison. The displayed example is an illustration, not a representative performance sample or a live model execution.

Comparator — direct	Difference	95% interval
One token	+0.0898	[−0.7483, 0.8449]
JSON	+160.2958	[138.1189, 177.0587]

Table 2: Primary paired HTTP latency differences in milliseconds. Intervals resample observed local family/tag clusters; JSON uses a different serialization prompt.

zero; there were no maximum-logit ties. The primary latency statistic is the median across 150 local items of the within-item difference between five-repetition mean latencies. It is not a difference between marginal medians.

Table 2 gives 10,000-replicate whole-family/tag percentile intervals. They describe sensitivity to the observed cluster composition, conditional on this model, machine, selected items, and session. They do not estimate between-machine or independent-session variability. The one-token interval crosses zero: this run provides no evidence of a material direct-readout latency advan-

tage. JSON was slower under its different serialization requirements, with task-dependent quality changes.

4.3 External task-specific results

Direct readout and one-token selection have identical external labels. JCoLA reached 86.0% accuracy and MCC 0.5708; JSON reached 82.0% and MCC 0.5160. The selected sample contains 158 acceptable sentences, so an always-acceptable rule already achieves 79.0% accuracy and has undefined MCC. Reporting only accuracy would obscure this imbalance.

JCommonsenseQA accuracy was 95.0% for direct/one-token and 95.5% for JSON, a difference of one example. JSTS retains its continuous reference scores in $[0, 5]$: direct expected-stage MAE at $T = 1$ was 0.5065. JSON has no comparable probability vector, so we do not assign it an expected-stage metric. The separate hard-stage comparison gives MAE 0.5480 for direct/one-token and 0.5650 for JSON. We do not round the gold scores

into invented classes or pool these tasks into one accuracy.

4.4 Presentation sensitivity and averaging

Two subsequent studies add Qwen2.5-0.5B-Instruct and Qwen2.5-1.5B-Instruct to the native 35B checkpoint. The first uses 7,600 requests on the same 600 external items. All 1,800 exact-input repeat pairs have identical logits and probabilities, yet display reversal changes native JCQA labels on 11/200 items. Conversely, the 1.5B checkpoint has zero JCoLA flips while always predicting acceptable: stability alone is insufficient.

The second freezes 600 additional public-development items before its 14,400 requests. It excludes prior IDs and literal text reuse, and selects JSTS pairs from previously unused sentence/image components. However, 194/200 JCoLA items share bibliographic sources with earlier evaluations. Both studies complete without inference errors; each excludes 57 synthetic checks/warm-ups.

We average semantic candidate probabilities at $T = 1$ across cyclic rotations of complete option lines, keeping output-token–meaning bindings fixed. Fixed-binding transformations and cyclic aggregation have precedent (Wei et al., 2024; Zheng et al., 2024); this is not a new algorithm or a reproduction of PriDe. Table 3 compares canonical-single and forward-cyclic accuracy, alongside single/reversed and forward/reverse-cyclic label disagreement. Lower disagreement does not consistently improve accuracy. The 0.5B cyclic JCoLA output is always acceptable (78.5% accuracy, undefined MCC). On JSTS, 0.5B and 1.5B cyclic hard labels occupy one stage on 199/200 and 200/200 items, respectively. These marginals expose collapse hidden by type validity or stability.

One cyclic answer costs 2, 5, or 6 model calls for JCoLA, JCQA, or JSTS; the single baseline reuses a pool member. Binary forward/reverse pools contain the same two calls, so their equality is structural. There is no equal-call repeat control in this averaging study. The small checkpoints use Transformers/MPS float32; the native model uses quantized Metal inference. Architecture, precision, and backend differ, precluding a model-size or cross-backend speed conclusion. These are descriptive system checks; fuller distributions, grouped intervals, and call accounting accompany the artifacts.

Checkpoint	Accuracy (%)	Flip rate (%)
Qwen2.5 0.5B	39.5 → 36.0	65.5 → 30.5
Qwen2.5 1.5B	80.0 → 83.5	26.0 → 12.5
Qwen3.6 35B-A3B	98.5 → 98.0	4.5 → 1.5

Table 3: JCQA on the additional 200-item panel: single → cyclic point estimates. Flip rate measures orientation disagreement, not error rate. Native 35B uses Q4_K_M. These are not equal-compute comparisons.

4.5 Reproduction and framework integration

The versioned source/evidence package at github.com/UpHash-Network/mini-jev records the three studies’ 26,050 requests, protocols, selections, hashes, analyzers, and AI-agent audits. Clean-directory CPU reanalysis reproduced all 56 derived files byte-for-byte. A fresh native build served all three types using a cached, hash-verified model on the same Mac. This validates reconstruction with source/model cache reuse, not a second-machine or fresh-download replication.

Separate synthetic diagnostics test integration boundaries. With the same checkpoint, complete prefixes and candidate IDs, unmodified LMQL 0.7.3 scoring through a custom native backend matched all 12 direct-readout labels. Maximum probability difference was 2.63×10^{-5} ; three cases exceeded the prespecified 10^{-5} tolerance. Different native scoring paths therefore did not establish exact numerical parity. ChainForge 0.3.7.4’s registration/dispatch routes, with a custom provider calling the released API, preserved all eight typed-answer pairs across four request pairs (maximum probability difference zero); three invalid inputs were rejected. This checks backend interoperability, not the browser flow editor, human usability, or task quality. These diagnostics are separate from the three evaluation panels. Inference replication still requires upstream data, pinned models, and compatible runtimes; hardware/build equivalence is not guaranteed.

5 Related Work and Scope

Candidate and ordinal scoring. Verbalizers (Schick and Schütze, 2021), decoding-free selection (Ma et al., 2025), Qwen3 reranking (Zhang et al., 2025; Qwen Team, 2025), and G-Eval’s weighted scores (Liu et al., 2023) precede our readouts. Zhuang et al. (2024) compute expected relevance from normalized fine-grained label log-likelihoods in zero-shot ranking. Wang

et al. (2026) use grade-token expectations for job-candidate ranking and fine-tune with a combined mean-squared-error and cross-entropy loss. Score shares this established expectation; frozen-weight scoring is also not new. Our contribution concerns typed inspection and system evaluation, not ranking gains or a new scoring rule.

Structured interfaces. LMQL exposes continuation scoring and selection (LMQL contributors, 2026); multi-token scoring need not equal our single-token readout. Outlines, Guidance, and grammar-constrained decoding support structured generation (dottxt and Outlines contributors, 2026; Guidance contributors, 2026; Geng et al., 2023). TypeSafe’s adapter maps general LLM outputs to a Jev-style contract (TypeSafe, 2026a); generated probability values differ from token-derived candidate probabilities. Our JSON experiment is a serialization comparison, not a framework benchmark.

Inspection interfaces. Rakutko’s distinct minijev compares frozen letter-logit scoring with constrained JSON in an inspection bench (Rakutko, 2026); our rename avoids confusion. Ollama documents local Choice/Noul/Score probabilities (Ollama, 2026); ChainForge adds an Ollama decision provider (Arawjo et al., 2024; ChainForge contributors, 2026). Its inference path remains untested here. ChainForge 0.3.7.6’s GUI retains six matched saved-evidence cases; its native Processor also executes our custom JavaScript to reconstruct 2,400 conditions from 6,000 archived logit records within 10^{-12} tolerance. This supports competitor feasibility, not inspection superiority (Appendix A). LLM Comparator supports response comparisons (Kahng et al., 2025); Langfuse exposes typed distributions (Langfuse, 2026), and Phoenix preserves Choice metadata (Arize AI and Phoenix contributors, 2026). Our distinction concerns connecting complete panels, semantic alignment, shared calls, and archive-line provenance in one workflow, not typed inspection alone.

Sensitivity and calibration. Order/token bias and fixed-binding changes are established (Wei et al., 2024; Zheng et al., 2024). Recent preprints distinguish reduced order sensitivity from accuracy (Hanna and Feng, 2026), and compressed ordinal hard-label distributions from per-item probability concentration in JEV/KEV (Gao et al., 2026). These overlap our stability and collapse

observations; our frozen-model cyclic study is not a reproduction of their experiments. Rao and Callison-Burch (2026) report correlated judge errors limiting confidence cascades in saved-verdict replay; we do not test such routing. We make these distinctions inspectable alongside measured call costs, without claiming them as new general findings. Temperature scaling (Guo et al., 2017) differs from contextual label-bias correction (Zhao et al., 2021); canonical sorting does not establish invariance to option meaning or placement.

Limitations

Presentation studies cover three checkpoints on one Apple Silicon device; matched parity/latency covers one native quantized model in one session. Architecture, precision, and backend differences prevent model-size attribution. Japanese public subsets have unknown training contamination and source dependence; the English interface example does not establish English task quality. Local data are AI-authored and development-related, without independent human reference validation. No user study establishes usability or downstream decision quality.

Candidate probabilities vary with prompts, labels, coverage, and temperature. Concentrated wrong answers occur; type validity and entropy concentration cannot replace application-specific validation before consequential actions. Score assumes equally spaced stages; Noul is a relative candidate probability, not a verified truth frequency. Questions run sequentially on a large model; other operating systems and lower-memory configurations remain unvalidated. Live-demo timings are not directly comparable with the research harness’s controlled interval.

Acknowledgements

OpenAI Codex agents assisted implementation, question authoring, checks, experiments, analysis, literature organization, and drafting under the author’s direction. Independent audit code was AI-assisted, not independent human review; AI-authored items lack human-expert validation. The author is responsible for all claims.

Ethics and Availability

The demonstration supports model inspection and typed-interface prototyping, with no evidence for high-stakes deployment. Inference is local; users

remain responsible for inputs and downstream use. Code and self-authored local data use MIT terms. External-dataset-derived selection meta-data, rubrics, protocols, and results separately retain applicable CC BY-SA 4.0 attribution and share-alike terms. Upstream dataset text, model weights, and compiled binaries are not bundled; repository data cards and notices specify acquisition and attribution.

References

Diogo Almeida. 2026. [Introducing system one models & jev](#). Published 2026-09-15; accessed 2026-09-20.

Ian Arawjo, Chelse Swoopes, Priyan Vaithilingam, Martin Wattenberg, and Elena L. Glassman. 2024. [ChainForge: A visual toolkit for prompt engineering and LLM hypothesis testing](#). In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, pages 1–18. Association for Computing Machinery.

Arize AI and Phoenix contributors. 2026. [Phoenix evaluators: Support for AI SDK evaluation models](#). Source change included in JavaScript phoenix-evals 2.6.0, released September 22, 2026; accessed September 27, 2026.

ChainForge contributors. 2026. [ChainForge: Ollama decision-model scoring](#). Source change of October 1, 2026; accessed October 2, 2026. Separate from the version 0.3.7.4 custom-provider diagnostic.

dottxt and Outlines contributors. 2026. [Outlines: Output types](#). Official documentation; accessed September 25, 2026.

Tianxiang Gao, Jinzhe Li, Zhiyuan Li, Yi Chang, and Yuan Wu. 2026. [More choices, fewer decisions: Ordinal-scale bias in JEV-like direct-decision models](#). *arXiv preprint arXiv:2609.38827*. Version 1, September 30, 2026.

Saibo Geng, Martin Josifoski, Maxime Peyrard, and Robert West. 2023. [Grammar-constrained decoding for structured NLP tasks without finetuning](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 10932–10952, Singapore. Association for Computational Linguistics.

ggml-org. 2026. [Qwen3.6-35b-a3b gguf](#). Pinned Q4_K_M artifact; accessed 2026-09-20.

ggml-org and contributors. 2026. [llama.cpp](#). Pinned revision used by the native helper.

Guidance contributors. 2026. [Guidance: A guidance language for controlling large language models](#). Official repository; accessed September 25, 2026.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. [On calibration of modern neural networks](#). In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR.

Karl Hanna and Chen Feng. 2026. [Accuracy and order sensitivity diverge under label-free strategies](#). *arXiv preprint arXiv:2608.11947*. Version 1, August 12, 2026; accessed September 25, 2026.

Minsuk Kahng, Ian Tenney, Mahima Pushkarna, Michael Xieyang Liu, James Wexler, Emily Reif, Krystal Kallarackal, Minsuk Chang, Michael Terry, and Lucas Dixon. 2025. [LLM Comparator: Interactive analysis of side-by-side evaluation of large language models](#). *IEEE Transactions on Visualization and Computer Graphics*, 31(1).

Kentaro Kurihara, Daisuke Kawahara, and Tomohide Shibata. 2022. [JGLUE: Japanese general language understanding evaluation](#). In *Proceedings of the Thirteenth Language Resources and Evaluation Conference*, pages 2957–2966. European Language Resources Association.

Langfuse. 2026. [Jev as a judge](#). Official documentation for the experimental evaluator; accessed September 27, 2026.

Yaniv Leviathan, Matan Kalman, and Yossi Matias. 2025. [Prompt repetition improves non-reasoning llms](#). *arXiv preprint arXiv:2512.14982*.

Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023. [G-eval: NLG evaluation using GPT-4 with better human alignment](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2511–2522, Singapore. Association for Computational Linguistics.

LMQL contributors. 2026. [LMQL generations API: Continuation scoring](#). Official documentation; accessed September 25, 2026.

Mingyu Derek Ma, Yanna Ding, Zijie Huang, Jianxi Gao, Yizhou Sun, and Wei Wang. 2025. [Inferring from logits: Exploring best practices for decoding-free generative candidate selection](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 33125–33144, Vienna, Austria. Association for Computational Linguistics.

Ollama. 2026. [Decision](#). <https://docs.ollama.com/capabilities/decision>. Official documentation; accessed 2026-10-02. Local runtime not independently tested.

Qwen Team. 2025. [Qwen3-reranker-8b: Official model card and inference example](#). Accessed 2026-09-20; primary source for yes/no candidate-logit extraction.

- Qwen Team. 2026. [Qwen3.6-35b-a3b: Official model card](#). Source-model revision 995ad96eacd98c81ed38be0c5b274b04031597b0; accessed 2026-09-20. 640
- Mikhail Rakutko. 2026. mini-Jev: Read the letter. <https://github.com/r-ms/mini-jev>. Software and experimental technical note. Accessed 2026-09-29. 641
- Delip Rao and Chris Callison-Burch. 2026. [JEV vs. LLMs as rubric judges: Cheaper, faster, and wrong in the same places](#). *arXiv preprint arXiv:2609.29769*. Version 2, September 28, 2026. 642
- Timo Schick and Hinrich Schütze. 2021. [Exploiting cloze-questions for few-shot text classification and natural language inference](#). In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 255–269. Association for Computational Linguistics. 643
- Taiga Someya, Yushi Sugimoto, and Yohei Oseki. 2024. [JCoLA: Japanese corpus of linguistic acceptability](#). In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation*, pages 9477–9488. ELRA and ICCL. 644
- TypeSafe. 2026a. [System one adapter for python](#). Official repository; accessed September 25, 2026. 645
- TypeSafe. 2026b. [Typesafe api reference](#). Accessed 2026-09-20; interface inspiration, not implementation equivalence. 646
- Qihang Wang, Jinwei Tan, Mengyuan Shi, Mayank Sharma, Shuai Zhao, Fuxian Li, Ryan Yan, Alexander P. Kreuzer, Mohit Jain, Dheeraj Toshniwal, and Manoj Seethamsetty. 2026. [Single-token expected-value scoring for cold-start candidate ranking](#). *arXiv preprint arXiv:2609.18188*. Version 1, September 16, 2026. 647
- Sheng-Lun Wei, Cheng-Kuang Wu, Hen-Hsen Huang, and Hsin-Hsi Chen. 2024. [Unveiling selection biases: Exploring order and token sensitivity in large language models](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 5598–5621, Bangkok, Thailand. Association for Computational Linguistics. 648
- Yanzhao Zhang, Mingxin Li, Dingkun Long, Xin Zhang, Huan Lin, Baosong Yang, Pengjun Xie, An Yang, Dayiheng Liu, Junyang Lin, Fei Huang, and Jingren Zhou. 2025. [Qwen3 embedding: Advancing text embedding and reranking through foundation models](#). *arXiv preprint arXiv:2506.05176*. Section 2 describes yes/no next-token reranking; inference details are also available in the official Qwen3-Reranker model card. 649
- Zihao Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. 2021. [Calibrate before use: Improving few-shot performance of language models](#). In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 12697–12706. PMLR. 650
- Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. 2024. [Large language models are not robust multiple choice selectors](#). In *The Twelfth International Conference on Learning Representations*. ICLR 2024 Spotlight; author manuscript version 4, February 22, 2024; conference record <https://openreview.net/forum?id=shr9PXz7T0>. 651
- Honglei Zhuang, Zhen Qin, Kai Hui, Junru Wu, Le Yan, Xuanhui Wang, and Michael Bendersky. 2024. [Beyond yes and no: Improving zero-shot LLM rankers via scoring fine-grained relevance labels](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 358–370, Mexico City, Mexico. Association for Computational Linguistics. 652

A Comparison Scope and Evidence Access

Comparison	Shared conditions	Recorded observation	Unmeasured or different
Native one-token	Checkpoint, full prefix, candidate IDs/logits, greedy selection	1,350/1,350 labels match; paired latency interval includes zero	Other machines and independent sessions
Grammar JSON	Resident model, selected tasks, HTTP timing interval	Validity, task-specific quality and serialization latency	Prompt and generated length differ; no probability-vector parity
LMQL 0.7.3	Checkpoint, full prefix, candidate IDs; custom native backend	12/12 labels match; 3/12 probability comparisons exceed 10^{-5}	Native evaluation paths differ; not a stock-backend benchmark
ChainForge 0.3.7.4	Same local API, request content and typed-response contract	Eight answer pairs identical; three invalid inputs rejected	Browser flow editor and human performance
ChainForge 0.3.7.6 GUI	Six saved cases; separate raw-logit/recipe input with custom JavaScript	Saved exports agree; native Processor reconstructs 2,400 conditions within 10^{-12}	Custom diagnostic code; no inference or human performance
Langfuse Phoenix	/ Primary-source feature descriptions only	Typed evaluation and probability inspection/metadata precedents	No matched inference, GUI trial or comparative user outcome

Table 4: Evidence tiers, not a capability ranking. Unmeasured properties are not absent capabilities. Integration diagnostics are separate from the three research panels.

Public and internal inspection. The released HTTP response and browser expose semantic labels, typed values, candidate probabilities, concentration, and probability-semantics metadata. Raw candidate token IDs and logits belong to the internal engine and research traces; they are not fields in the public response. Table 4 distinguishes controlled comparisons, integration checks, and source-based related work. These observations do not establish that people work faster or make better decisions with this interface.

Model-free evidence review. The linked repository includes a frozen source/evidence ZIP and a reviewer guide. Its standard-library verification checks the archive and recorded request accounting without weights, network access, or inference. A separate CPU replay executes the six frozen analyses in a new directory and compares all 56 derived files. The three studies contain 26,050 measured requests; 300 earlier pilot predictions are retained as background and excluded from that count. The current manuscript and later framework diagnostics are separate from the frozen ZIP. Replaying recorded data verifies calculations, not the truth of the original observations or their generalization.

The explorer is a later static view of the two presentation studies, not an alteration of their frozen records. Its exporter verifies the archive, internal hashes, completion inventories, physical candidate mappings, and saved aggregation recipes. Original

numeric values remain unchanged. The underlying cyclic and dihedral variants average semantic probabilities, not logits. Binary forward/reverse cyclic pools share their two physical calls, so identical outputs there are structural. Software checks cover every exported record, comparison arithmetic, shared-call accounting, and JSON export round trips; they are not human task results.

Live installation. The repository supplies native build and pinned-model acquisition commands: `build_native.sh`, `fetch_native_model.py`, and `./run.sh -without-calibration`. Running `python3 -m demo.server -port 8766` serves the interface at <http://127.0.0.1:8766/>. Fresh builds use $T = 1$; historical calibration artifacts are bound to their recorded runtime. The tested Mac uses macOS 26.4, Metal, and Accelerate. The 20.4 GB model requires additional runtime memory; this is not a demonstrated minimum-memory requirement. We distribute source, not model weights or compiled binaries. The static explorer can instead be served from `docs/` with Python’s HTTP server; HTTPS or localhost enables browser hash verification.

Post-hoc LMQL diagnosis. Recomputing the same 12 saved comparisons preserves the three 10^{-5} probability-tolerance failures (score-01, score-03, score-04). Maximum Score expectation error is 7.18×10^{-5} , within the separate, original 10^{-4} typed-value tolerance. All native top-two candidate probability margins ex-

ceed 0.9799; matching labels therefore provide no near-boundary test. Candidate softmax over saved LMTP logits reproduces saved LMQL probabilities within 3.64×10^{-9} . Recorded native-path logit differences can account for the observed discrepancies; their underlying computational cause remains unisolated. Type, candidate count, and assistant prefix are confounded, so the failures do not identify a Score-specific defect. This analysis adds no inference calls, changes no threshold, and does not convert the failed probability comparisons into passes.

Matched saved-evidence workflow. The `comparison package` freezes six ID-hash-selected cases across two studies and three types before GUI execution, with 32 diagnostic question instances. Prior results were known; this is not blinded preregistration. All presenters receive the complete selected panel and source index. The unmodified published ChainForge wheel imports adapted response caches; native table search, response lightboxes, spreadsheet export, and full flow export retain the supplied evidence. The adapter supplies named probabilities, alignment, deltas, and call unions; generic JSON/export access is not a native solution to every diagnostic question. No new model calls, human task scores, or efficiency claims follow. Setup, scripts, failures of exploratory automation, and exact outputs accompany the protocol.

Reconstruction from raw records. A `follow-up package` supplies only archived logits, candidate mappings, temperatures, gold references, and member recipes for those six panels. The unmodified ChainForge GUI executes supplied JavaScript from an empty output cache in Firefox 153. It recomputes softmax, semantic averaging, reference differences, shared-call unions, and panel summaries. An independently coded Python oracle checks the export: 2,400 original conditions agree with unchanged saved values; maximum numeric difference from JavaScript is 1.78×10^{-15} . Six tuple-reversal controls preserve results and eight malformed inputs are rejected. These finite software checks add no model observations or human evidence. The competitor needs custom code; our explorer uses preprocessed data and does not ingest this raw fixture. Neither implementation establishes a measured practical advantage.

Human evaluation boundary. No human study was conducted. An earlier synthetic-page rehearsal is material QA, excluded from experimental results. A counterbalanced pilot packet now targets the released explorer and complete-JSON view, with no participants or observations. Human comparative performance and cross-machine replication remain open.